

Relational Model

- Table = relation.
- Column headers = *attributes*.
- Row = *tuple*

| name | manf |
|------------|--------|
| WinterBrew | Pete's |
| BudLite | A.B. |
| ... | ... |

Beers

- *Relation schema* = name(attributes) + other structure info., e.g., keys, other constraints.
Example: Beers(name, manf).
 - ◆ Order of attributes is arbitrary, but in practice we need to assume the order given in the relation schema.
- *Relation instance* is current set of rows for a relation schema.
- *Database schema* = collection of relation schemas.

Why Relations?


- Very simple model.
- *Often* a good match for the way we think about our data.
- Abstract model that underlies SQL, the most important language in DBMS's today.
 - ◆ But SQL uses “bags,” while the abstract relational model is set-oriented.

Relational Design

Simplest approach (not always best): convert each E.S. to a relation and each relationship to a relation.

Entity Set \rightarrow Relation

E.S. attributes become relational attributes.




Becomes:

`Beers(name, manf)`

E/R Relationships \rightarrow Relations

Relation has attribute for *key* attributes of each E.S. that participates in the relationship.

- Add any attributes that belong to the relationship itself.
- Renaming attributes OK.
 - ◆ Essential if multiple roles for an E.S.



Likes(drinker, beer)
 Favorite(drinker, beer)
 Buddies(name1, name2)
 Married(husband, wife)

Combining Relations

Sometimes it makes sense to combine relations.

- Common case: Relation for an E.S. E plus the relation for some many-one relationship from E to another E.S.

Example

Combine `Drinker(name, addr)` with `Favorite(drinker, beer)` to get `Drinker1(name, addr, favBeer)`.

- Danger in pushing this idea too far: redundancy.
- e.g., combining `Drinker` with `Likes` causes the drinker's address to be repeated viz.:


| name | addr | beer |
|-------|-----------|--------|
| Sally | 123 Maple | Bud |
| Sally | 123 Maple | Miller |

- Notice the difference: `Favorite` is many-one; `Likes` is many-many.

Weak Entity Sets, Relationships \rightarrow Relations

- Relation for a weak E.S. must include its full key (i.e., attributes of related entity sets) as well as its own attributes.
- A supporting (double-diamond) relationship yields a relation that is actually redundant and should be deleted from the database schema.

Example



Hosts(hostName)

Logins(loginName, hostName)

At(loginName, hostName, hostName2)


- In At, hostName and hostName2 must be the same host, so delete one of them.
- Then, Logins and At become the same relation; delete one of them.
- In this case, Hosts' schema is a subset of Logins' schema. Delete Hosts?

Subclasses → Relations

Three approaches:

1. Object-oriented: each entity is in one class.
Create a relation for each class, with all the attributes for that class.
 - ◆ Don't forget inherited attributes.
2. E/R style: an entity is in a network of classes related by **isa**. Create one relation for each E.S.
 - ◆ An entity is represented in the relation for each subclass to which it belongs.
 - ◆ Relation has only the attributes attached to that E.S. + key.
3. Use nulls. Create one relation for the root class or root E.S., with all attributes found anywhere in its network of subclasses.
 - ◆ Put NULL in attributes not relevant to a given entity.

Example



OO-Style

| name | manf |
|------|------|
| Bud | A.B. |

Beers

| name | manf | color |
|------------|--------|-------|
| SummerBrew | Pete's | dark |

Ales

E/R Style

| name | manf |
|------------|--------|
| Bud | A.B. |
| SummerBrew | Pete's |

Beers

| name | color |
|------------|-------|
| SummerBrew | dark |

Ales

Using Nulls

| name | manf | color |
|------------|--------|-------|
| Bud | A.B. | NULL |
| SummerBrew | Pete's | dark |

Beers

Functional Dependencies

$X \rightarrow A$ = assertion about a relation R that whenever two tuples agree on all the attributes of X , then they must also agree on attribute A .

Example

Drinkers(name, addr, beersLiked, manf, favoriteBeer)

| name | addr | beersLiked | manf | favoriteBeer |
|---------|------------|------------|--------|--------------|
| Janeway | Voyager | Bud | A.B. | WickedAle |
| Janeway | Voyager | WickedAle | Pete's | WickedAle |
| Spock | Enterprise | Bud | A.B. | Bud |

- Reasonable FD's to assert:
 1. name \rightarrow addr
 2. name \rightarrow favoriteBeer
 3. beersLiked \rightarrow manf

- Shorthand: combine FD's with common left side by concatenating their right sides.
- Sometimes, several attributes jointly determine another attribute, although neither does by itself. Example:

`beer bar` \rightarrow `price`

Keys of Relations

K is a *key* for relation R if:

1. $K \rightarrow$ all attributes of R .
 2. For **no proper subset** of K is (1) true.
- If K at least satisfies (1), then K is a *superkey*.

Conventions

- Pick one key; underline key attributes in the relation schema.
- X , etc., represent sets of attributes; A etc., represent single attributes.
- No set formers in FD's, e.g., ABC instead of $\{A, B, C\}$.

Example

Drinkers(name, addr, beersLiked, manf, favoriteBeer)

- {name, beersLiked} FD's all attributes, as seen.
 - ◆ Shows {name, beersLiked} is a superkey.
- name \rightarrow beersLiked is false, so name not a superkey.
- beersLiked \rightarrow name also false, so beersLiked not a superkey.
- Thus, {name, beersLiked} is a key.
- No other keys in this example.
 - ◆ Neither name nor beersLiked is on the right of any observed FD, so they must be part of *any* superkey.
- Important point: “key” in a relation refers to tuples, not the entities they represent. If an entity is represented by several tuples, then entity-key will not be the same as relation-key.

Who Determines Keys/FD's?

- We could assert a key K .
 - ◆ Then the only FD's asserted are that $K \rightarrow A$ for every attribute A .
 - ◆ No surprise: K is then the only key for those FD's, according to the formal definition of “key.”
- Or, we could assert some FD's and *deduce* one or more keys by the formal definition.
 - ◆ E/R diagram implies FD's by key declarations and many-one relationship declarations.
- Rule of thumb: FD's either come from keyness, many-1 relationship, or from physics.
 - ◆ E.g., “no two courses can meet in the same room at the same time” yields $\text{room time} \rightarrow \text{course}$.

Inferring FD's

And this is important because . . .

- When we talk about improving relational designs, we often need to ask “does this FD hold in this relation?”


Given FD's $X_1 \rightarrow A_1, X_2 \rightarrow A_2 \dots X_n \rightarrow A_n$, does FD $Y \rightarrow B$ necessarily hold in the same relation?

- Start by assuming two tuples agree in Y . Use given FD's to infer other attributes on which they must agree. If B is among them, then yes, else no.

Algorithm

Define $Y^+ = \text{closure of } Y = \text{set of attributes functionally determined by } Y$:

- Basis: $Y^+ := Y$.
- Induction: If $X \subseteq Y^+$, and $X \rightarrow A$ is a given FD, then add A to Y^+ .




- End when Y^+ cannot be changed.

Example

$A \rightarrow B, BC \rightarrow D.$

- $A^+ = AB.$
- $C^+ = C.$
- $(AC)^+ = ABCD.$



Finding All Implied FD's

Motivation: Suppose we have a relation $ABCD$ with some FD's F . If we decide to decompose $ABCD$ into ABC and AD , what are the FD's for ABC , AD ?

- Example: $F = AB \rightarrow C, C \rightarrow D, D \rightarrow A$.
It looks like just $AB \rightarrow C$ holds in ABC , but in fact $C \rightarrow A$ follows from F and applies to relation ABC .
- Problem is exponential in worst case.

Algorithm

- For each set of attributes X compute X^+ .
 - ◆ But skip $X = \emptyset$, $X = \text{all attributes}$.
 - ◆ Add $X \rightarrow A$ for each A in $X^+ - X$.
- Drop $XY \rightarrow A$ if $X \rightarrow A$ holds.
- Finally, project the FD's by selecting only those FD's that involve only the attributes of the projection.
 - ◆ Notice that after we project the discovered FD's onto some relation, the eliminated FD's can be inferred *in the projected relation*.

Example

In ABC with FD's $A \rightarrow B$, $B \rightarrow C$, project onto AC .

1. $A^+ = ABC$; yields $A \rightarrow B$, $A \rightarrow C$.
 2. $B^+ = BC$; yields $B \rightarrow C$.
 3. $AB^+ = ABC$; yields $AB \rightarrow C$; drop in favor of $A \rightarrow C$.
 4. $AC^+ = ABC$ yields $AC \rightarrow B$; drop in favor of $A \rightarrow B$.
 5. $C^+ = C$ and $BC^+ = BC$; adds nothing.
- Resulting FD's: $A \rightarrow B$, $A \rightarrow C$, $B \rightarrow C$.
 - Projection onto AC : $A \rightarrow C$.